



Logistic Regression

Simple logistic regression




$$z = \theta_0 + \theta_1 \cdot x$$

Multivariate Logistic Regression




$$z = \theta_0 + \theta_1 \cdot x_1 + \theta_2 \cdot x_2$$


Which one is a better classifier?



P1




P2





P3

Multivariate Logistic Regression

(3 dimensional example)







Logistic regression steps:


- ▶ Select a random θ as the initial parameter
- ▶ Calculate $\hat{y} = \sigma(\theta^T x)$
- ▶ Calculate the difference between \hat{y} and actual y

$$\text{cost}(\hat{y}, y) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$$


- ▶ Calculate cost function for all training samples ($J(\theta)$)
- ▶ Update the parameters by gradient descent

$$\theta_j = \theta_{j-1} - \gamma \nabla J(\theta)$$


Which is the best Learning rate



Small learning rate



Large Learning Rate



Small Learning Rate

Evaluating the Logistic Regression Model

- ▶ **Accuracy**
- ▶ **Precision**
- ▶ **Recall**
- ▶ **F1 score**
- ▶ **Confusion Matrix**

Sklearn implimentation

```
class sklearn.linear_model.LogisticRegression(penalty='l2', *, dual=False,  
tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1, class_weight=None,  
random_state=None, solver='lbfgs', max_iter=100, multi_class='deprecated',  
verbose=0, warm_start=False, n_jobs=None, l1_ratio=None) \[source\]
```

Logistic Regression (aka logit, MaxEnt) classifier.

In the multiclass case, the training algorithm uses the one-vs-rest (OvR) scheme if the 'multi_class' option is set to 'ovr', and uses the cross-entropy loss if the 'multi_class' option is set to 'multinomial'. (Currently the 'multinomial' option is supported only by the 'lbfgs', 'sag', 'saga' and 'newton-cg' solvers.)

This class implements regularized logistic regression using the 'liblinear' library, 'newton-cg', 'sag', 'saga' and 'lbfgs' solvers. **Note that regularization is applied by default.** It can handle both dense and sparse input. Use C-ordered arrays or CSR matrices containing 64-bit floats for optimal performance; any other input format will be converted (and copied).

The 'newton-cg', 'sag', and 'lbfgs' solvers support only L2 regularization with primal formulation, or no regularization. The 'liblinear' solver supports both L1 and L2 regularization, with a dual formulation only for the L2 penalty. The Elastic-Net regularization is only supported by the 'saga' solver.

Read more in the [User Guide](#).

penalty : {'l1', 'l2', 'elasticnet', None}, default='l2'

Specify the norm of the penalty:

- **None** : no penalty is added;
- **'l2'** : add a L2 penalty term and it is the default choice;
- **'l1'** : add a L1 penalty term;
- **'elasticnet'** : both L1 and L2 penalty terms are added.

penalty	$r(w)$	
None	0	
ℓ_1	$\ w\ _1$	Lasso cost function
ℓ_2	$\frac{1}{2}\ w\ _2^2 = \frac{1}{2}w^T w$	Ridge cost function
ElasticNet	$\frac{1-\rho}{2}w^T w + \rho\ w\ _1$	

`solver : {'lbfgs', 'liblinear', 'newton-cg', 'newton-cholesky', 'sag', 'saga'}, default='lbfgs'`

Algorithm to use in the optimization problem. Default is 'lbfgs'. To choose a solver, you might want to consider the following aspects:

- For small datasets, 'liblinear' is a good choice, whereas 'sag' and 'saga' are faster for large ones;
- For multiclass problems, only 'newton-cg', 'sag', 'saga' and 'lbfgs' handle multinomial loss;
- 'liblinear' and 'newton-cholesky' can only handle binary classification by default. To apply a one-versus-rest scheme for the multiclass setting one can wrap it with the `OneVsRestClassifier`.
- 'newton-cholesky' is a good choice for `n_samples >> n_features`, especially with one-hot encoded categorical features with rare categories. Be aware that the memory usage of this solver has a quadratic dependency on `n_features` because it explicitly computes the Hessian matrix.